

# Generative modeling of mixed tabular data with the R package `arf`

Jan Kapar<sup>1,2</sup>, David S. Watson, Kristin Blesch, and Marvin N. Wright

<sup>1</sup>Leibniz Institute for Prevention Research & Epidemiology – BIPS

<sup>2</sup>Faculty of Mathematics and Computer Science, University of Bremen

July 10, 2024

UseR! 2024

# Generative modeling of mixed tabular data with arf

Motivation

---



2

The hype about generative modeling

# Generative modeling of mixed tabular data with arf

## Motivation

---



2

## The hype about generative modeling



Generative AI is revolutionizing content creation by producing human-like text, images, and music. Its ability to generate high-quality, original content quickly is driving widespread interest and adoption.

# Generative modeling of mixed tabular data with arf

## Motivation

---



2

## The hype about generative modeling



Generative AI is revolutionizing content creation by producing human-like text, images, and music. Its ability to generate high-quality, original content quickly is driving widespread interest and adoption.

# Generative modeling of mixed tabular data with arf

Motivation



2

The hype about generative modeling



# Generative modeling of mixed tabular data with arf

## Motivation

---



2

## The hype about generative modeling

- Text, speech, image and video synthesis
- Deep-learning-based methods
- Implementations almost exclusively in Python

# Generative modeling of mixed tabular data with arf

## Motivation

---



2

## The hype about generative modeling

- Text, speech, image and video synthesis
- Deep-learning-based methods
- Implementations almost exclusively in Python

## What about tabular data?

- Less research
- Existing deep learning solutions often fail

# Generative modeling of mixed tabular data with `arf`

## Motivation

---



2

## The hype about generative modeling

- Text, speech, image and video synthesis
- Deep-learning-based methods
- Implementations almost exclusively in Python

## What about tabular data?

- Less research
- Existing deep learning solutions often fail

⇒ `arf` package for tabular data synthesis based on `ranger`



# Generative modeling

## Definition

---



3

Given a data set consisting of vectors  $x \in \mathbb{R}^d$ :

# Generative modeling

## Definition

---



3

Given a data set consisting of vectors  $\boldsymbol{x} \in \mathbb{R}^d$ :

- **Generative** modeling = modeling the **joint density**  $p(\boldsymbol{x})$

# Generative modeling

## Definition

---



3

Given a data set consisting of vectors  $\mathbf{x} \in \mathbb{R}^d$ :

- **Generative** modeling = modeling the **joint density**  $p(\mathbf{x})$   
(vs. **discriminative** modeling = directly modeling the **conditional density**  $p(y|X = \mathbf{x})$  for some dependent variable  $y \in \mathbb{R}$ )

# Generative modeling

## Definition

---



3

Given a data set consisting of vectors  $\boldsymbol{x} \in \mathbb{R}^d$ :

- **Generative** modeling = modeling the **joint density**  $p(\boldsymbol{x})$   
(vs. **discriminative** modeling = directly modeling the **conditional density**  $p(y|X = \boldsymbol{x})$  for some dependent variable  $y \in \mathbb{R}$ )
- Not always explicit density estimation

# Generative modeling

## Use cases for synthetic data

---



- Direct usage:
  - Chatbots
  - Translation
  - Transcription
  - Image generation
  - ...

# Generative modeling

## Use cases for synthetic data

---



4

- Direct usage:
  - Chatbots
  - Translation
  - Transcription
  - Image generation
  - ...
- Indirect usage:
  - Data **augmentation**
  - Data **balancing**
  - Missing data **imputation**
  - **Privacy** preservation
  - ...

# The package `arf`

## Overview

---



5

Implementation of adversarial random forests<sup>1</sup>:

---

<sup>1</sup>Watson et al. [2023]

# The package arf

## Overview



5

## Implementation of adversarial random forests<sup>1</sup>:

- Fast
  - Built on `data.table`
  - Optional parallelization via `future`, `doParallel`
  - Fast random forest implementation `ranger`



---

<sup>1</sup>Watson et al. [2023]



# The package `arf`

## Overview

---



5

## Implementation of adversarial random forests<sup>1</sup>:

- Fast
  - Built on `data.table`
  - Optional parallelization via `future`, `doParallel`
  - Fast random forest implementation `ranger`
- Easy-to-use



---

<sup>1</sup>Watson et al. [2023]

# The package `arf`

## Overview

---

### Implementation of adversarial random forests<sup>1</sup>:

- Fast
  - Built on `data.table`
  - Optional parallelization via `future`, `doParallel`
  - Fast random forest implementation `ranger`
- Easy-to-use
- No tuning required



---

<sup>1</sup>Watson et al. [2023]

# The package `arf`

## Overview

### Implementation of adversarial random forests<sup>1</sup>:

- Fast
  - Built on `data.table`
  - Optional parallelization via `future`, `doParallel`
  - Fast random forest implementation `ranger`
- Easy-to-use
- No tuning required
- Able to handle missing values (`ranger` branch 'missing\_values')



---

<sup>1</sup>Watson et al. [2023]

# The package `arf`

## Overview



5

## Implementation of adversarial random forests<sup>1</sup>:

- Fast
  - Built on `data.table`
  - Optional parallelization via `future`, `doParallel`
  - Fast random forest implementation `ranger`
- Easy-to-use
- No tuning required
- Able to handle missing values (`ranger` branch 'missing\_values')
- Comparable or superior performance on (mixed) tabular data



---

<sup>1</sup>Watson et al. [2023]

# The package `arf`

Algorithm and functions

---



6

Step 1: Grow adversarial random forests (`adversarial_rf()`)

# The package `arf`

## Algorithm and functions

---



6

### Step 1: Grow adversarial random forests (`adversarial_rf()`)

- Iterative approach: Train random forest classifiers to distinguish real data from naive synthetic data drawn from marginals of real data

# The package arf

## Algorithm and functions



6

### Step 1: Grow adversarial random forests (`adversarial_rf()`)

- Iterative approach: Train random forest classifiers to distinguish real data from naive synthetic data drawn from marginals of real data
- Partitioning of data manifold into areas where **independence w.r.t. the variables** can be assumed locally

# The package `arf`

## Algorithm and functions

---



6

### Step 1: Grow adversarial random forests (`adversarial_rf()`)

- Iterative approach: Train random forest classifiers to distinguish real data from naive synthetic data drawn from marginals of real data
- Partitioning of data manifold into areas where **independence w.r.t. the variables** can be assumed locally

### Step 2: Estimate mixture distribution (`forde()`)



# The package `arf`

## Algorithm and functions



6

### Step 1: Grow adversarial random forests (`adversarial_rf()`)

- Iterative approach: Train random forest classifiers to distinguish real data from naive synthetic data drawn from marginals of real data
- Partitioning of data manifold into areas where **independence w.r.t. the variables** can be assumed locally

### Step 2: Estimate mixture distribution (`forde()`)

- Estimate univariate densities per variable within leaves

# The package `arf`

## Algorithm and functions



6

### Step 1: Grow adversarial random forests (`adversarial_rf()`)

- Iterative approach: Train random forest classifiers to distinguish real data from naive synthetic data drawn from marginals of real data
- Partitioning of data manifold into areas where **independence w.r.t. the variables** can be assumed locally

### Step 2: Estimate mixture distribution (`forde()`)

- Estimate univariate densities per variable within leaves
- Form mixture distribution over leaves weighted by share of real data

# The package `arf`

Algorithm and functions

---



Step 3: Use mixture distribution for

# The package `arf`

Algorithm and functions

---



7

Step 3: Use mixture distribution for

- (Conditional) sampling (`forge()`)

# The package `arf`

Algorithm and functions

---



7

## Step 3: Use mixture distribution for

- (Conditional) sampling (`forge()`)
- (Conditional) likelihood computation (`lik()`)

# The package `arf`

## Algorithm and functions

---



### Step 3: Use mixture distribution for

- (Conditional) sampling (`forge()`)
- (Conditional) likelihood computation (`lik()`)
- (Conditional) expectation computation (`expct()`)

# The package arf

## Machine learning utility evaluation



8

Dataset	Characteristics	Model	Accuracy (sd)	F1 (sd)	Time in sec
adult	classes = 2 $n = 32,561$ $d = 14$	Real	0.828 (0.006)	0.884 (0.004)	
		ARF	<b>0.819</b> (0.006)	<b>0.877</b> (0.005)	<b>2.9</b>
		CTGAN	0.786 (0.020)	0.853 (0.019)	263.3
		CTABGAN+	0.808 (0.008)	0.869 (0.006)	561.6
		TVAE	0.804 (0.007)	0.865 (0.006)	115.1
census	classes = 2 $n = 298,006$ $d = 40$	Real	0.922 (0.002)	0.957 (0.001)	
		ARF	0.903 (0.019)	0.946 (0.012)	<b>53.2</b>
		CTGAN	0.916 (0.015)	0.954 (0.009)	4287.8
		CTABGAN+	0.912 (0.026)	0.952 (0.016)	10182.1
		TVAE	<b>0.928</b> (0.007)	<b>0.961</b> (0.004)	1814.9
covertypes	classes = 7 $n = 581,012$ $d = 54$	Real	0.895 (0.000)	0.838 (0.000)	
		ARF	<b>0.707</b> (0.006)	<b>0.549</b> (0.006)	<b>103.5</b>
		CTGAN	0.633 (0.009)	0.400 (0.009)	13387.2
		CTABGAN+	NA	NA	>24h
		TVAE	0.698 (0.013)	0.459 (0.013)	4882.0

# The package arf

## Machine learning utility evaluation



9

Dataset	Characteristics	Model	Accuracy (sd)	F1 (sd)	Time in sec
credit	classes = 2 $n = 284,807$ $d = 30$	<i>Real</i>	<i>0.997 (0.001)</i>	<i>0.607 (0.029)</i>	
		ARF	0.995 (0.001)	<b>0.527</b> (0.036)	<b>32.2</b>
		CTGAN	0.881 (0.099)	0.047 (0.031)	4898.0
		CTABGAN+	<b>0.998</b> (0.000)	0.000 (0.000)	7497.3
		TVAE	<b>0.998</b> (0.000)	0.000 (0.000)	3847.6
intrusion	classes = 5 $n = 494,021$ $d = 40$	<i>Real</i>	<i>0.998 (0.001)</i>	<i>0.833 (0.001)</i>	
		ARF	<b>0.993</b> (0.001)	<b>0.656</b> (0.001)	<b>68.2</b>
		CTGAN	0.944 (0.088)	0.645 (0.088)	8749.3
		CTABGAN+	NA	NA	>24h
		TVAE	0.990 (0.002)	0.598 (0.002)	4306.0

Omitted IT-GAN and RCC-GAN.




# The package arf

## Runtime



10



# The package `arf`

Simple examples: Data synthesis

---



11

Generate synthetic samples from the `palmerpenguins` dataset:

# The package arf

## Simple examples: Data synthesis



11

Generate synthetic samples from the palmerpenguins dataset:

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
forge(params, 5)
```

```
## # A tibble: 5 x 8
##   species    island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>      <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie    Dream           41.4           17.1             197           3380
## 2 Adelie    Biscoe           42            18.8             190           4152
## 3 Chinstrap Dream           51.3           19.2             193           3407
## 4 Adelie    Dream           36.5           17.2             191           3779
## 5 Adelie    Biscoe           40.5           19.2             193           3919
## # i 2 more variables: sex <fct>, year <int>
```

# The package arf

## Simple examples: Data synthesis



11

Generate synthetic samples from the palmerpenguins dataset:

```
library(arf)
arf <- adversarial_rf(penguins, num_trees = 10, min_node_size = 2)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
forge(params, 5)
```

```
## # A tibble: 5 x 8
##   species    island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>      <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie    Dream           41.4           17.1             197           3380
## 2 Adelie    Biscoe           42             18.8             190           4152
## 3 Chinstrap Dream           51.3           19.2             193           3407
## 4 Adelie    Dream           36.5           17.2             191           3779
## 5 Adelie    Biscoe           40.5           19.2             193           3919
## # i 2 more variables: sex <fct>, year <int>
```

# The package arf

## Simple examples: Data synthesis



11

Generate synthetic samples from the palmerpenguins dataset:

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins, family = 'truncnorm', finite_bounds = 'no')
forge(params, 5)
```

```
## # A tibble: 5 x 8
##   species    island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>      <fct>         <dbl>         <dbl>           <int>       <int>
## 1 Adelie    Dream           41.4           17.1             197        3380
## 2 Adelie    Biscoe           42            18.8             190        4152
## 3 Chinstrap Dream           51.3           19.2             193        3407
## 4 Adelie    Dream           36.5           17.2             191        3779
## 5 Adelie    Biscoe           40.5           19.2             193        3919
## # i 2 more variables: sex <fct>, year <int>
```

# The package `arf`

Simple examples: Data synthesis

---



12

Generate samples under some desired conditions:

# The package arf

## Simple examples: Data synthesis



12

Generate samples under some desired conditions:

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
forge(params, 5, evidence = data.frame(island = 'Torgersen'))
```

```
## # A tibble: 5 x 8
##   species island    bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>          <dbl>          <dbl>          <int>        <int>
## 1 Adelie Torgersen      45.8           20.8           192         4612
## 2 Adelie Torgersen      45.6           19.3           196         3484
## 3 Adelie Torgersen      36            20.9           190         3746
## 4 Adelie Torgersen      36.6           20.8           191         3563
## 5 Adelie Torgersen      40.6           17.7           180         3162
## # i 2 more variables: sex <fct>, year <int>
```

# The package arf

## Simple examples: Data synthesis



12

Generate samples under some desired conditions:

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
forge(params, 5, evidence = data.frame(island = 'Torgersen', bill_length_mm = '>40'))
```

```
## # A tibble: 5 x 8
##   species island    bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>          <dbl>         <dbl>         <int>         <int>
## 1 Adelie Torgersen      41.7           19.8           208           3647
## 2 Adelie Torgersen      41.1           21.3           200           4535
## 3 Adelie Torgersen      45.4           18.8           204           4852
## 4 Adelie Torgersen      40.2           19.3           197           3661
## 5 Adelie Torgersen      43.6           20.9           194           4268
## # i 2 more variables: sex <fct>, year <int>
```



# The package `arf`

Simple examples: (Conditional) likelihoods

---



# The package arf

## Simple examples: (Conditional) likelihoods

13

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
lik(params, penguins[1:3,])
```

```
## [1] -3.655597 -4.023504 -4.154264
```

# The package arf

## Simple examples: (Conditional) likelihoods

13

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
lik(params, penguins[1:3,], log = F)
```

```
## [1] 0.02584605 0.01789017 0.01569734
```

# The package arf

## Simple examples: (Conditional) likelihoods



13

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
lik(params, penguins[1:3,], evidence = data.frame(sex = 'female'))
```

```
## [1] -4.179873 -3.654580 -4.259066
```

# The package `arf`

Simple examples: (Conditional) expectations

---



# The package arf

## Simple examples: (Conditional) expectations



14

```
library(arf)
arf <- adversarial_rf(penguins)
```

```
## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%
```

```
params <- forde(arf, penguins)
expct(params)
```

```
## # A tibble: 1 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>   <fct>         <dbl>         <dbl>         <dbl>         <dbl>
## 1 Adelie  Biscoe          43.9           17.2           201.          4200.
## # i 2 more variables: sex <fct>, year <int>
```

# The package arf

## Simple examples: (Conditional) expectations



14

```
library(arf)
arf <- adversarial_rf(penguins)

## Iteration: 0, Accuracy: 79.88%
## Iteration: 1, Accuracy: 44.95%

params <- forde(arf, penguins)
expct(params, query = "flipper_length_mm", evidence = data.frame(body_mass_g = 3300))

## # A tibble: 1 x 1
##   flipper_length_mm
##               <dbl>
## 1               189.
```

# The package `arf`

Simple examples: Shortcut functions

---





# The package arf

## Simple examples: Shortcut functions



15

```
library(arf)
rarf(penguins, 3)
```

```
## # A tibble: 3 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>    <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie   Dream           41.1           17.6             200          3408
## 2 Adelie   Biscoe           41.4            19             183          3749
## 3 Chinstrap Dream           50.7           19.4             190          3484
## # i 2 more variables: sex <fct>, year <int>
```

```
darf(penguins, query = penguins[1:3,])
```

```
## [1] -5.490055 -5.394553 -5.905513
```

```
earf(penguins)
```

```
## # A tibble: 1 x 8
##   species island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>    <fct>         <dbl>         <dbl>           <dbl>         <dbl>
## 1 Adelie   Biscoe           43.9           17.2             201.          4201.
```

# The package `arf`

Simple examples: Parallelization

---



# The package arf

## Simple examples: Parallelization



16

```
library(arf)
doFuture::registerDoFuture()
future::plan("multisession", workers = 4)
rarf(penguins, 5)
```

```
## # A tibble: 5 x 8
##   species  island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>    <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Chinstrap Dream          52.1           17.9             185           3409
## 2 Gentoo   Biscoe          44.3           13.9             210           5479
## 3 Chinstrap Dream          51.5           19.8             198           4063
## 4 Gentoo   Biscoe          50.7           14.9             219           6256
## 5 Gentoo   Biscoe          49.3           16.1             224           5951
## # i 2 more variables: sex <fct>, year <int>
```

# The package arf

## Simple examples: Parallelization



16

```
library(arf)
doParallel::registerDoParallel(cores = 4)

rarf(penguins, 5)
```

```
## # A tibble: 5 x 8
##   species    island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>      <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie    Dream             41.4           17.1             197           3380
## 2 Adelie    Biscoe             42            18.8             190           4152
## 3 Chinstrap Dream             51.3           19.2             193           3407
## 4 Adelie    Dream             36.5           17.2             191           3779
## 5 Adelie    Biscoe             40.5           19.2             193           3919
## # i 2 more variables: sex <fct>, year <int>
```

# The package arf

## Simple examples: Parallelization



16

```
library(arf)
doParallel::registerDoParallel(cores = 4)
```

```
rarf(penguins, 5, parallel = F)
```

```
## # A tibble: 5 x 8
##   species    island bill_length_mm bill_depth_mm flipper_length_mm body_mass_g
##   <fct>      <fct>         <dbl>         <dbl>           <int>         <int>
## 1 Adelie    Dream             41.4           17.1             197           3380
## 2 Adelie    Biscoe             42            18.8             190           4152
## 3 Chinstrap Dream             51.3           19.2             193           3407
## 4 Adelie    Dream             36.5           17.2             191           3779
## 5 Adelie    Biscoe             40.5           19.2             193           3919
## # i 2 more variables: sex <fct>, year <int>
```

# Application

Statistical utility: Setup

---



17

Replication of published German national cohort health studies with  
arf-generated synthetic data.

# Application

Statistical utility: Setup

---



17

Replication of published German national cohort health studies with arf-generated synthetic data.

Procedure:

- Learn 5 arf models

# Application

## Statistical utility: Setup

---



17

Replication of published German national cohort health studies with arf-generated synthetic data.

Procedure:

- Learn 5 arf models
  - Sample 5 data sets from each
- ⇒ 25 synthetic datasets



# Application

## Statistical utility: Setup

---



17

Replication of published German national cohort health studies with `arf`-generated synthetic data.

Procedure:

- Learn 5 `arf` models
  - Sample 5 data sets from each
- ⇒ 25 synthetic datasets
- Perform study analysis on every synthetic dataset
  - Calculate median and 95% confidence intervals

# Application

## Statistical utility: Setup



17

Replication of published German national cohort health studies with `arf`-generated synthetic data.

Procedure:

- Learn 5 `arf` models
  - Sample 5 data sets from each
- ⇒ 25 synthetic datasets
- Perform study analysis on every synthetic dataset
  - Calculate median and 95% confidence intervals
  - Compare with original data results

# Application

Wienbergen et al. [2022] ( $n = 1,713$ ,  $d = 12$ ) - Myocardial infarction

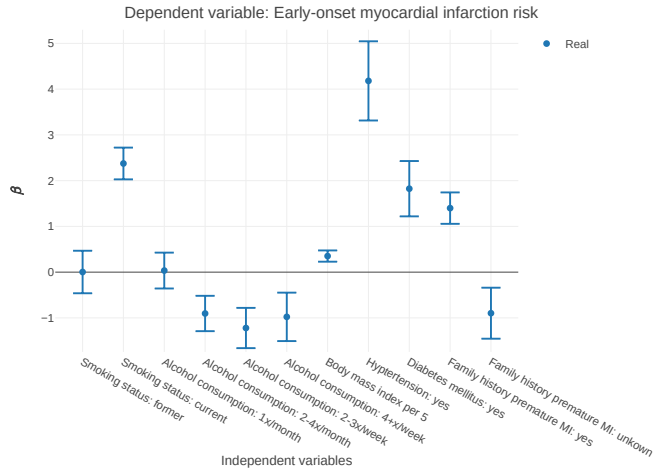
---



# Application

Wienbergen et al. [2022] ( $n = 1,713$ ,  $d = 12$ ) - Myocardial infarction


Logistic regression:



# Application

Wienbergen et al. [2022] ( $n = 1,713$ ,  $d = 12$ ) - Myocardial infarction

Logistic regression:



# Application

Berger et al. [2021] ( $n = 82,205$ ,  $d = 47$ ) - Loneliness during COVID

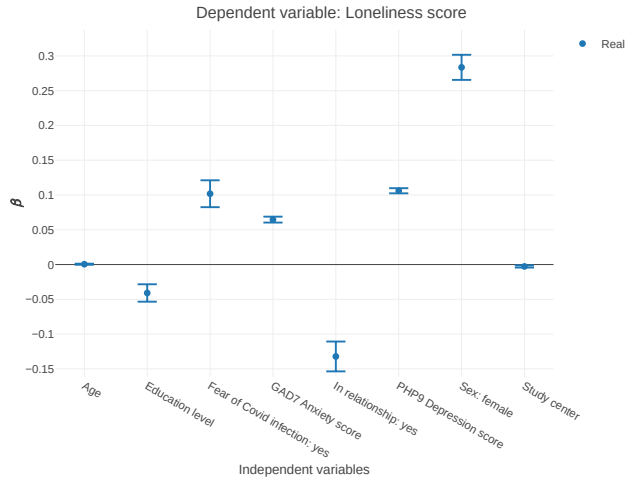
---



# Application

Berger et al. [2021] ( $n = 82,205$ ,  $d = 47$ ) - Loneliness during COVID


Linear regression:



# Application

Berger et al. [2021] ( $n = 82,205$ ,  $d = 47$ ) - Loneliness during COVID

Linear regression:





# Conclusion & Outlook

---



20

## The package `arf`

- Little or no tuning required
- Fast computation
- Competitive performance on tabular data

# Conclusion & Outlook

---

## The package `arf`

- Little or no tuning required
- Fast computation
- Competitive performance on tabular data

## Outlook

- (Conditional) variances, covariances, CDFs
- Missing data imputation
- Explainable AI
- Counterfactuals generation
- Privacy guarantees: Differential privacy

## Download arf



<https://cran.r-project.org/package=arf>  
<https://github.com/bips-hb/arf>

- K. Berger, S. Riedel-Heller, A. Pabst, M. Rietschel, and D. Richter. Loneliness during the first wave of the SARS-CoV-2 pandemic—results of the German National Cohort (NAKO). [Bundesgesundheitsblatt-Gesundheitsforschung-Gesundheitsschutz](#), 64:1157–1164, 2021.
- D. S. Watson, K. Blesch, J. Kapar, and M. N. Wright. Adversarial random forests for density estimation and generative modeling. In [International Conference on Artificial Intelligence and Statistics](#), pages 5357–5375. PMLR, 2023.
- H. Wienbergen, D. Boakye, K. Günther, J. Schmucker, L. A. Mata Marín, H. Kerniss, R. Nagrani, L. Struß, S. Rühle, T. Retzlaff, et al. Lifestyle and metabolic risk factors in patients with early-onset myocardial infarction: a case-control study. [European Journal of Preventive Cardiology](#), 29(16):2076–2087, 2022.
- L. Xu, M. Skoularidou, A. Cuesta-Infante, and K. Veeramachaneni. Modeling tabular data using conditional GAN. In [Advances in Neural Information Processing Systems](#), volume 32, 2019.

# Thank you for your attention

[www.leibniz-bips.de/en](http://www.leibniz-bips.de/en)

## Contact

Jan Kapar

Leibniz Institute for Prevention Research  
and Epidemiology – BIPS

Achterstraße 30  
D-28359 Bremen

[kapar@leibniz-bips.de](mailto:kapar@leibniz-bips.de)



# The package arf

## Machine learning utility evaluation (Setup)

